

Crucible

Governed agentic software creation.

A governed, agentic SDLC — specialized agents, quality gates between every phase, one shared record. From idea to ship: governance, auditability, transparency.

Ember Agentic Labs · Tony Gibbons, founder · tony@emberagenticlabs.com

Pre-seed · 2026

AI coding tools give you code that runs. That's where they stop.

Point a coding agent at real work and you get something that works in the moment — but undocumented, untested, inconsistent across a growing codebase, and drifting further the longer the project runs.


Past a prototype, “it runs” isn’t “it ships.” The gap isn’t the model’s capability — it’s everything that has to be true around the code.

56%

of CEOs report no cost or revenue gains from AI

12%

report both
(PwC 2026, n=4,454)




Coherence decays as the work grows — the failure a bigger model doesn't fix.

Deterministic gates keep a mercurial foundation model honest.

Foundation models are mercurial by nature, and no orchestration layer changes that — an LLM will sometimes skip a step, lose the thread, or claim something is done when it isn't.


The design response is a deterministic boundary around a probabilistic actor: state machines and gates that read the shared record and decide what is allowed to happen next. Before a tool call leaves an agent, the system asks whether there is evidence the tests actually ran before this “complete” call.

A check either passes or it does not — it is not a judgment call the model can argue with.



An invalid handoff is blocked, and the agent self-corrects — live, in the product.

Structure built in as the work is created — not cleaned up after.



Checks at the point of action

A bad step is refused, not flagged after. Specialized agents, each conditioned to its role, not one generalist.

The output is the whole SDLC

Plans, requirements, architecture, tests and docs — governed artifacts, not just code.

One shared record

Every artifact is a typed node joined by typed edges. Tracing a bug back to the requirement that drove it is a graph query, not a memory exercise.

Local, on your own keys

A VS Code extension. Signed installers for macOS, Windows and Linux. Shipped May 2026 — no waitlist.

You direct; specialized agents pick up the work. Nothing becomes an artifact until it clears the gate.

Users pay the model provider directly, so revenue carries no compute cost.

Alpha — free

Runs on the user's own keys. Costless to try; drives bottom-up developer adoption.

Pro — ~\$20-30/mo

Deeper orchestration, extended autonomy, multi-model.

Enterprise — high ARPU

Audit-ready artifacts, enforced governance, full traceability, local operation.

≈100% gross margin

No compute COGS, and near-zero customer-acquisition cost — unusual for a developer tool, and the reason the free alpha costs almost nothing to distribute.

A better foundation model is a more capable individual; the lasting advantage is the organization around it. So Crucible compounds as models improve, instead of being replaced by them.

**Every software-producing team is a user;
the paying market splits in two. Both figures are potential at
scale — Crucible is pre-revenue today.**

\$30M

Potential ARR — Pro at 100,000 seats × \$25/month, from a professional-developer population in the tens of millions. Illustrative, not a forecast.
professional-developer population in the tens of millions.

\$100M

Potential ARR — Enterprise at 1,000 organisations, midpoint of the \$50–150K/year governance-software band. Illustrative, not a forecast.
the \$50–150K/year governance-software band.

The pool it comes out of

\$6.37T

worldwide IT spend forecast for 2026 (Gartner, July 2026)

\$1.52T

US accumulated software technical debt (CISQ, 2022)

17 hrs

per developer per week on maintenance, debugging and refactoring
(Stripe / Harris Poll, 2018)

*The market has begun pricing this specific layer: 8090
Labs raised a \$135M Series A (Salesforce Ventures) in June
2026 to attack this same layer.*

The moat is the accumulated governance — the rules earned from real use.

A fast-follower can copy the architecture. What they cannot copy is the year of “the model skipped X, so we now enforce X” — rules earned from real use, accumulated in the shared record and applied by the gates on every task from then on. That balance grows with use and cannot be shortcut.

Execution tools are swappable backends, not rivals. The comparison that matters is not which tool writes code faster, but which produces work an organisation will actually accept.

Coding agents

Cursor · Claude Code · Copilot · Aider

Execution engines — they write the code, fast. Crucible runs on top of them as swappable backends. Not rivals.

Orchestration platforms

Devin · MetaGPT · Factory · Augment · 8090 Labs

The real comparison, and several are further along and better-resourced — but they are enterprise cloud. Crucible is complete, governed, free and local.

The difference is architectural: the orchestration platforms treat governance as review after the fact. Crucible enforces it at the point of creation — hard to retrofit onto a single-autonomous-agent design.

Distribution, not engineering, is the constraint.

1

Bottom-up first

The free local alpha costs a developer nothing to try, so adoption compounds without a sales motion.

2

A GTM cofounder

A governance layer is a non-obvious sell. The near-term priority is a cofounder to own the category narrative and the commercial motion.

3

Regulated design partners

Life sciences, defence and finance — where local operation is the unlock and willingness to pay is highest. The first named design partner is a deep-tech operating company.

Why governance sells into regulated work

In life sciences, defence and finance the alternative to auditable AI output isn't a competitor — it's not being allowed to use AI on the work at all.

Audit-ready artifacts, enforced gates and requirement-to-test traceability are the buying criteria, and local operation means regulated data never leaves the building.

The pre-seed and the cofounder search exist to lift exactly that constraint.

The alpha shipped publicly on 2026-05-28, and it is already in production use.

The first named user is a deep-tech operating company. They evaluated Crucible against their own work and against competitive products, adopted it on that basis, and run the alpha.

Pre-revenue by design: the free alpha ships with telemetry off by default, so there's no active-user count — and I haven't marketed it.

Full-time since January 2026. Crucible built most of Crucible, including generating its own product website from its own project record — every task through the same gates.



January 2026

Full-time on Crucible.



28 May 2026

Alpha ships publicly — signed installers for macOS, Windows and Linux. Free, local, no waitlist.



Since the alpha shipped

A deep-tech operating company evaluates Crucible against their own work and against competitive products — and adopts it. First named user.

Shipped, installable today, and in use — before a dollar of marketing.

The cash cost was ~\$600. The real capital was several months of senior engineering time.

~\$600

total model spend to build the alpha, over about three months

\$325/mo

current burn — franchise tax, code-signing, domains, hosting, tooling, model spend

\$0

salary taken — full-time since January 2026.
A senior technology leader's time, at no cash cost

Near-zero cash burn is real, but it is not free — the company has been funded by the foregone earnings of a senior technology leader. That is why the raise funds go-to-market and a cofounder, not engineering.

Bottom-up projections

Pro \$25/month × 100,000 seats

\$30M potential ARR

Enterprise \$50–150K/year × 1,000 organisations (midpoint)

\$100M potential ARR

A few hundred million ARR at software margins in a good outcome, enterprise-weighted.

Two decades building systems where correctness is enforced before the fact.

Tony Gibbons — solo founder

Intel

ASIC design — correctness enforced before fabrication.

Google Cloud, ten years

Built the unified revenue platform, 0 → 8,000 daily users. Ran the \$10B+ CapEx long-range plan. Led an ML program that saved \$100M in a quarter. Scaled the technology-partnership org from 50 to 225 as chief of staff.

UC Berkeley

Master of Information and Data Science, 2025.

Independently evaluated

An independent evaluation at a deep-tech operating company, with the tools driven by a UC Berkeley AI researcher who holds a PhD and founded a major bank's AI lab.

Actively recruiting

A go-to-market cofounder to own the category narrative and the commercial motion.

Bootstrapped. \$0 raised, 100% founder-owned.

No outside investment taken

Ember Agentic Labs is a California LLC.

Preparing a ~\$1M pre-seed

To fund go-to-market and the first regulated-industry deployments.

No SBIR or NIH funding

And actively recruiting a go-to-market cofounder.